

Super Market Sales Analytics

1) Data Science Proposal

1. The Team

Group Number: 49

Members:

- Binsu Elizabeth Varghese(2021C104187)
- Shehza Fathima (2021C104174)
- Nithin Krishnan (2021C104176)

2. Problem Statement and Background

To Analyze and Build models for the sales of different products in a supermarket located in different cities using machine learning techniques

3. The Data Source

The dataset is Supermarket_Sales_Dataset.csv. This dataset contains various attributes of supermarkets like Invoice Id, Branch, City, Customer Type, Gender, Product Type, Unit Price, Quantity, Tax, Selling Price, Date, Time, Payment Type, Cost Price, Gross Income, Rating.

4. Goals of Your Analysis

- Get maximum insights from a data set
- Uncover Underlying structure
- Extract important features from the data set
- Train a Machine Learning model for predicting Customer Rating
- Validation of Predicted Model
- Visualization of results with Graphical representations

5. Description of Data Analysis Tools You Plan to Use

- numpy
- pandas

- seaborn
- matplotlib
- sklearn
- scipy

6. Describe the Data Products Your Project Will Produce

- Performance of 3 clustering techniques - AHC, K-Means, K-Medoids
- Comparison of performance of the two classifiers – Logistic regression and Decision tree to predict

In [1]:

```
1  #Import Libraries
2  import numpy as np
3  import pandas as pd
4  import seaborn as sns
5  from matplotlib import pyplot as plt
6  from scipy.stats import skew
7  from prettytable import PrettyTable
8  from datetime import datetime
9
10 # For feature selections and Feature Engineering
11 from sklearn.model_selection import train_test_split
12 from sklearn.feature_selection import SelectKBest, chi2, mutual_info_classif
13 from sklearn.preprocessing import MinMaxScaler
14 from sklearn.preprocessing import LabelEncoder
15 from sklearn.preprocessing import StandardScaler
16 from mlxtend.feature_selection import SequentialFeatureSelector as SFS
17 from sklearn.linear_model import LinearRegression
18 from collections import Counter
19 from sklearn.decomposition import PCA
20 from imblearn.combine import SMOTETomek
21 from sklearn.ensemble import RandomForestClassifier
22
23 # Clustering
24 from sklearn.cluster import KMeans
25 from sklearn_extra.cluster import KMedoids
26 from sklearn.cluster import AgglomerativeClustering
27 from sklearn.metrics import silhouette_score
28 import silhouetteplot
29 import scipy.cluster.hierarchy as sch
30 from sklearn.datasets import make_moons
31 from sklearn.metrics import adjusted_rand_score
32
33 # Importing the DecisionTreeClassifier and LogisticRegressionClassifier for model building
34 from sklearn import tree
35 from sklearn.linear_model import LogisticRegression
36 from sklearn.tree import DecisionTreeClassifier
37 from sklearn import model_selection
38
39 # For Analyzing the models
40 from sklearn.metrics import confusion_matrix
41 from sklearn.metrics import f1_score
```

```

42 from sklearn.metrics import roc_curve, roc_auc_score
43 from sklearn.metrics import (accuracy_score, classification_report, confusion_matrix)
44 from sklearn.model_selection import cross_val_score
45
46 import warnings
47 warnings.simplefilter(action="ignore", category=FutureWarning)
48 pd.options.mode.chained_assignment = None

```

In [2]:

```

1 #Read Data from the csv file
2 df = pd.read_csv('Supermarket_Sales_Dataset.csv')
3 df.head(2)

```

Out[2]:

	Invoice ID	Branch	City	Customer Type	Gender	Product Type	Unit Price	Quantity	Tax	Selling Price	Date	Time	Payment Type	Cost Price	Gross Income	Rating
0	750-67-8428	A	Bangalore	Member	Female	Health and beauty	74.69	7	26.1415	548.9715	1/5/2019	13:08	Debit card	522.83	26.1415	9.0
1	226-31-3081	C	Mysore	Normal	Female	Electronic accessories	15.28	5	3.8200	80.2200	3/8/2019	10:29	Cash	76.40	3.8200	9.0

In [3]:

```


1 #Defining class colour
2 class color:
3     BLUE = '\033[94m'
4     BOLD = '\033[1m'
5     END = '\033[0m'

```

Objectives

1) Which city has a better sale for products in the Electronic Accessories product line?

```
In [4]: 1 dfCat=df[['City','Product Type','Gross Income']].groupby(['City','Product Type'],as_index=False).sum()
2 sns.catplot(data=dfCat, y='Gross Income', x='City',kind='bar', hue='Product Type',height=5,aspect=2,palette="Set3")
3 plt.show()
4
5 dfCat[dfCat['Product Type'] == 'Electronic accessories'].nlargest(1,'Gross Income')
```



Out[4]:

	City	Product Type	Gross Income
12	Mysore	Electronic accessories	903.2845

2) Which payment method is used more often at a particular city, branch and for which product type ?

```
In [5]: 1 dfpm=df[['City','Product Type','Payment Type','Invoice ID']].groupby(['City','Product Type','Payment Type'],as_index=False)
2 dfpm.sort_values(by = ['City', 'Invoice ID'], ascending = [True, False],inplace = True)
3 dfpm.drop_duplicates(subset="City",inplace = True)
4 dfpm.head()
```

Out[5]:

	City	Product Type	Payment Type	Invoice ID
14	Bangalore	Home and lifestyle	Debit card	26
33	Mangalore	Sports and travel	Cash	26
42	Mysore	Food and beverages	Cash	31

3) Which Product type has been more purchased by female customers?

```
In [6]: 1 dfgender=df[['Gender','Product Type','Quantity']].groupby(['Gender','Product Type'],as_index=False).sum()
2 dfgender[dfgender['Gender'] == 'Female'].nlargest(1,'Quantity')
```

Out[6]:

	Gender	Product Type	Quantity
1	Female	Fashion accessories	530

4) In which month does the highest number of home and lifestyle products have been sold ?

```
In [7]: 1 dfm = df[df['Product Type'] == 'Home and lifestyle']
2 dfm['Date'] = pd.to_datetime(dfm['Date'])
3 dfm['Month'] = dfm['Date'].dt.month_name(locale = 'English')
4 dfmtrend = dfm[['Month', 'Quantity']].groupby([dfm['Month']]).sum()
5 dfmtrend.sort_values(by = ['Quantity'], ascending = [False], inplace = True)
6 dfmtrend.head(1)
```

Out[7]:

	Quantity
Month	
March	364

5) At what time most of the female customers are purchasing products ?

```
In [8]: 1 dff = df[df['Gender'] == 'Female']
2 dfft = dff[['Gender', 'Time', 'Invoice ID']]
3 dfft['Invoice ID Count'] = dfft['Invoice ID']
4 dfff=dfft[['Gender', 'Time', 'Invoice ID Count']].groupby(['Gender', 'Time'],as_index=False).count()
5 dfff.sort_values(by = ['Invoice ID Count'], ascending = False, inplace = True)
6 dfff.drop_duplicates(subset="Time",inplace = True)
7 dfff.head(1)
```

Out[8]:

	Gender	Time	Invoice ID Count
168	Female	14:42	6

2) Exploratory Data Analysis

```
In [9]: 1 #Size of the Dataset
2 print(color.BLUE + color.BOLD + "\nSize of Dataset:" + color.END)
3 print(df.shape)
```

Size of Dataset:
(1000, 16)

```
In [10]: 1 # Attribute and its datatype
2 ptbl = PrettyTable()
3
4 for attribute in df.columns:
5     ptbl.field_names = ["Attribute Name", "Data Type"]
6     ptbl.add_row([attribute, df[attribute].dtype])
7
8 print(ptbl)
```

Attribute Name	Data Type
Invoice ID	object
Branch	object
City	object
Customer Type	object
Gender	object
Product Type	object
Unit Price	float64
Quantity	int64
Tax	float64
Selling Price	float64
Date	object
Time	object
Payment Type	object
Cost Price	float64
Gross Income	float64
Rating	float64


```
In [11]: 1 #Distribution Of Data
2 distTxt = color.BLUE + color.BOLD + "Data Distribution" + color.END
3 print(distTxt.center(100))
4 df.describe()
```

Data Distribution

Out[11]:

	Unit Price	Quantity	Tax	Selling Price	Cost Price	Gross Income	Rating
count	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000	1000.000000
mean	55.672130	5.510000	15.379369	322.966749	307.58738	15.379369	6.97270
std	26.494628	2.923431	11.708825	245.885335	234.17651	11.708825	1.71858
min	10.080000	1.000000	0.508500	10.678500	10.17000	0.508500	4.00000
25%	32.875000	3.000000	5.924875	124.422375	118.49750	5.924875	5.50000
50%	55.230000	5.000000	12.088000	253.848000	241.76000	12.088000	7.00000
75%	77.935000	8.000000	22.445250	471.350250	448.90500	22.445250	8.50000
max	99.960000	10.000000	49.650000	1042.650000	993.00000	49.650000	10.00000

```
In [12]: 1 columnsList = list(df)
2 categoricallist = list(df.select_dtypes(include=['object']).columns)
3 numericalList = list(set(columnsList) - set(categoricallist))
4
5 #Character or Numerical Data
6 print(color.BLUE + color.BOLD + "Categorical Data:" + color.END)
7 print(categoricallist)
8 print(color.BLUE + color.BOLD + "\nNumerical Data:" + color.END)
9 print(numericalList)
```

Categorical Data:

```
['Invoice ID', 'Branch', 'City', 'Customer Type', 'Gender', 'Product Type', 'Date', 'Time', 'Payment Type']
```

Numerical Data:

```
['Tax ', 'Quantity', 'Selling Price', 'Cost Price', 'Gross Income', 'Unit Price', 'Rating']
```

```
In [13]: 1 df['RatingLevel'] = pd.cut(x=df['Rating'],
2                               bins=[1, 6, 10],
3                               labels=['Low', 'High'])
4
5 #Balanced or Imbalanced Dataset
6 print(color.BLUE + color.BOLD + "Classification of Rating Level:" + color.END)
7 count = df['RatingLevel'].value_counts()
8 print(count)
9 count.plot.pie(autopct='%.2f')
```


Classification of Rating Level:

High 659

Low 341

Name: RatingLevel, dtype: int64

Out[13]: <AxesSubplot:ylabel='RatingLevel'>







Data Visualization






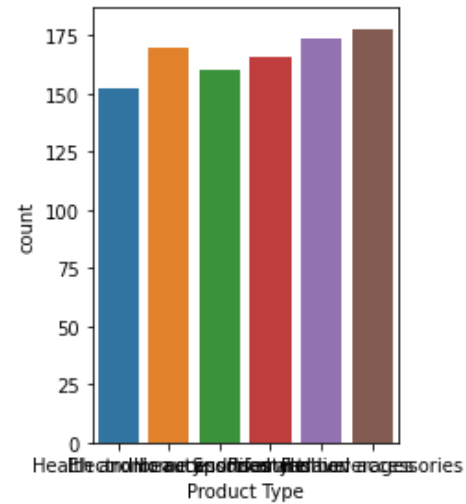

In [14]:

```

1  #Count Plots (for categorical attributes)
2
3  plt.rcParams["figure.figsize"] = [14.00, 4.0]
4  plt.rcParams["figure.autolayout"] = True
5
6  index = 0
7  graphsInARow = 4
8
9  for attr in categoricalList:
10
11     if (index % graphsInARow == 0):
12         f, ax = plt.subplots(1, graphsInARow)
13
14         sns.countplot(x=attr, data=df, ax = ax[index % graphsInARow])
15         index = index + 1
16
17     if (index % graphsInARow == 0):
18         plt.show()

```







In [15]:

```

1  #Histogram Distribution (for Continuous Attributes)
2  plt.rcParams["figure.figsize"] = [14.00, 4.0]
3  plt.rcParams["figure.autolayout"] = True
4
5  index = 0
6  graphsInARow = 4
7
8  for attr in numericalList:
9      if (index % graphsInARow == 0):
10         f, ax = plt.subplots(1, graphsInARow)
11         sns.histplot(data=df[attr], bins=10, kde=True, ax = ax[index % graphsInARow])
12         index = index + 1
13         if (index % graphsInARow == 0):
14             plt.show()

```






```
In [16]: 1 #Correlation of the features in the dataset
2 corr = df.corr()
3
4 corrTxt = color.BLUE + color.BOLD + "Correlation Using HeatMap" + color.END
5 print(corrTxt.center(120))
6
7 #Visualizing correlation using Heatmap
8 plt.figure(figsize=(10,7.5))
9 sns.heatmap(corr, annot=True, cmap='GnBu')
```

Correlation Using HeatMap

Out[16]: <AxesSubplot:>



3) Data Wrangling

```
In [17]: 1 #Count of NaN/Null values from dataset
2 print(color.BLUE + color.BOLD + "\nCount of NaN/Null values for each feature:" + color.END)
3 print(df.isna().sum())
```

Count of NaN/Null values for each feature:

```
Invoice ID      0
Branch          0
City            0
Customer Type   0
Gender          0
Product Type    0
Unit Price      0
Quantity        0
Tax             0
Selling Price   0
Date            0
Time            0
Payment Type    0
Cost Price      0
Gross Income    0
Rating          0
RatingLevel     0
dtype: int64
```

```
In [18]: 1 # Checking for duplicates
2 df.duplicated().sum()
```

Out[18]: 0

```
In [19]: 1 # Dropping the attributes that has a unique number(number assignment) for all the rows or attributes derivable from
2 df=df.drop(['Invoice ID','Branch', 'Date', 'Time'], axis=1)
```

```
In [20]: 1 df.shape
```

Out[20]: (1000, 13)

In [21]: 1 df.head(2)

Out[21]:

	City	Customer Type	Gender	Product Type	Unit Price	Quantity	Tax	Selling Price	Payment Type	Cost Price	Gross Income	Rating	RatingLevel
0	Bangalore	Member	Female	Health and beauty	74.69	7	26.1415	548.9715	Debit card	522.83	26.1415	9.1	High
1	Mysore	Normal	Female	Electronic accessories	15.28	5	3.8200	80.2200	Cash	76.40	3.8200	9.6	High

4) Clustering

```
In [22]: 1 #Copy of df
2 dfCopy = df.copy()
3
4 for column in dfCopy.columns:
5     #If Column data type is int or float continue
6     if dfCopy[column].dtype == 'int64' or dfCopy[column].dtype == 'float64':
7         continue
8     #If Column data type is object, encode and transform it
9     dfCopy[column] = LabelEncoder().fit_transform(dfCopy[column].astype(str))
```

In [23]: 1 dfCopy.head(2)

Out[23]:

	City	Customer Type	Gender	Product Type	Unit Price	Quantity	Tax	Selling Price	Payment Type	Cost Price	Gross Income	Rating	RatingLevel
0	0	0	0	3	74.69	7	26.1415	548.9715	2	522.83	26.1415	9.1	0
1	2	1	0	0	15.28	5	3.8200	80.2200	0	76.40	3.8200	9.6	0

```
In [24]: 1 # Scaling is done
          2 scaler = StandardScaler()
          3 scaled_features = scaler.fit_transform(dfCopy)
          4 scaled_features.shape
```

Out[24]: (1000, 13)

```
In [25]: 1 scaled_features_df = pd.DataFrame(scaled_features)
          2 scaled_features_df.columns = ['City', 'Customer Type', 'Gender', 'Product Type', 'Unit Price',
          3                               'Quantity', 'Tax ', 'Selling Price', 'Payment Type', 'Cost Price',
          4                               'Gross Income', 'Rating', 'RatingLevel']
          5 scaled_features_df.head()
```

Out[25]:

	City	Customer Type	Gender	Product Type	Unit Price	Quantity	Tax	Selling Price	Payment Type	Cost Price	Gross Income	Rating	RatingLevel
0	-1.208970	-0.998002	-0.998002	0.319617	0.718160	0.509930	0.919607	0.919607	1.203528	0.919607	0.919607	1.238443	-0.719340
1	1.238338	1.002002	-0.998002	-1.430109	-1.525303	-0.174540	-0.987730	-0.987730	-1.205937	-0.987730	-0.987730	1.529527	-0.719340
2	-1.208970	1.002002	1.002002	0.902859	-0.352781	0.509930	0.071446	0.071446	-0.001205	0.071446	0.071446	0.248760	-0.719340
3	-1.208970	-0.998002	1.002002	0.319617	0.096214	0.852165	0.675780	0.675780	1.203528	0.675780	0.675780	0.830927	-0.719340
4	-1.208970	1.002002	1.002002	1.486101	1.156959	0.509930	1.267125	1.267125	1.203528	1.267125	1.267125	-0.973790	1.390162

1. K-Means Clustering

```
In [26]: 1 range_n_clusters = [2, 3, 4]
2 silhouette_avg_mean = []
3 for num_clusters in range_n_clusters:
4
5     # initialise kmeans
6     kmeans1 = KMeans(n_clusters = num_clusters, max_iter = 50)
7     kmeans1.fit(dfCopy)
8
9     cluster_labels = kmeans1.labels_
10
11     # silhouette score
12     silhouette_avg = silhouette_score(dfCopy, cluster_labels).round(3)
13     silhouette_avg_mean.append(silhouette_avg)
14     print("For n_clusters={0}, the silhouette score is {1}".format(num_clusters, silhouette_avg))
```

For n_clusters=2, the silhouette score is 0.644

For n_clusters=3, the silhouette score is 0.591

For n_clusters=4, the silhouette score is 0.548

2. K-Medoids Clustering

```
In [27]: 1 range_n_clusters = [2, 3, 4]
2 silhouette_avg_med = []
3 for num_clusters in range_n_clusters:
4
5     kmedoids1 = KMedoids(num_clusters).fit(dfCopy)
6     cluster_labels = kmedoids1.labels_
7
8     # silhouette score
9     silhouette_avg = silhouette_score(dfCopy, cluster_labels).round(3)
10    silhouette_avg_med.append(silhouette_avg)
11    print("For n_clusters={0}, the silhouette score is {1}".format(num_clusters, silhouette_avg))
```

For n_clusters=2, the silhouette score is 0.631

For n_clusters=3, the silhouette score is 0.582

For n_clusters=4, the silhouette score is 0.535

3. Agglomerative Hierarchical Clustering (AHC)


```
In [28]: 1 range_n_clusters = [2, 3, 4]
2 silhouette_avg_clust = []
3 for num_clusters in range_n_clusters:
4     cluster1 = AgglomerativeClustering(n_clusters = num_clusters, affinity='euclidean', linkage='ward')
5     cluster1.fit_predict(dfCopy)
6     cluster_labels = cluster1.labels_
7
8     # silhouette score
9     silhouette_avg = silhouette_score(dfCopy, cluster_labels).round(3)
10    silhouette_avg_clust.append(silhouette_avg)
11    print("For n_clusters={0}, the silhouette score is {1}".format(num_clusters, silhouette_avg))
```

For n_clusters=2, the silhouette score is 0.639

For n_clusters=3, the silhouette score is 0.59


For n_clusters=4, the silhouette score is 0.518

```
In [29]: 1 plt.figure(figsize=(7, 5))
2
3 dendrogram = sch.dendrogram(sch.linkage(dfCopy, method = "ward"))
4 plt.title('Dendrogram')
5 plt.xlabel('Rating')
6 plt.ylabel('Euclidean distances')
7 plt.show()
```



Comparison of 3 Clustering methods using Elbow Method

```
In [30]: 1 plt.figure(figsize=(10, 5))
2
3 plt.style.use("fivethirtyeight")
4 plt.subplot(1,3,1)
5 plt.plot(range(2, 5), silhouette_avg_mean)
6 plt.xticks(range(2, 5))
7 plt.xlabel("Number of Clusters")
8 plt.ylabel("Silhouette Coefficient K-Means")
9
10 plt.style.use("fivethirtyeight")
11 plt.subplot(1,3,2)
12 plt.plot(range(2, 5), silhouette_avg_med)
13 plt.xticks(range(2, 5))
14 plt.xlabel("Number of Clusters")
15 plt.ylabel("Silhouette Coefficient K-Medoid")
16
17 plt.style.use("fivethirtyeight")
18 plt.subplot(1,3,3)
19 plt.plot(range(2, 5), silhouette_avg_clust)
20 plt.xticks(range(2, 5))
21 plt.xlabel("Number of Clusters")
22 plt.ylabel("Silhouette Coefficient AHC")
23
24 plt.subplots_adjust(left=1,
25                     bottom=0.1,
26                     right=2,
27                     top=0.9,
28                     wspace=1,
29                     hspace=1)
30 plt.show()
```


Best Silhouette Score

```
In [31]: 1 features, true_labels = make_moons(  
2         n_samples=250, noise=0.05, random_state=42  
3     )  
4 scaled_features1 = scaler.fit_transform(features)
```

In [32]:

```
1  # Best score is for k = 2
2  kmeans1 = KMeans(n_clusters=2)
3  kmedoids1 = KMedoids(n_clusters=2)
4  ahc1 = AgglomerativeClustering(n_clusters=2)
5
6  # Fit the algorithms to the features
7  kmeans1.fit(scaled_features1)
8  kmedoids1.fit(scaled_features1)
9  ahc1.fit_predict(scaled_features1)
10
11 # Compute the silhouette scores for each algorithm
12 kmeans_silhouette1 = silhouette_score(
13     scaled_features1, kmeans1.labels_).round(3)
14 kmedoids_silhouette1 = silhouette_score(
15     scaled_features1, kmedoids1.labels_).round (3)
16 ahc_silhouette1 = silhouette_score(
17     scaled_features1, ahc1.labels_).round (3)
```

```
In [33]: 1 # Plot the data and cluster silhouette comparison
2 fig, (ax1, ax2) = plt.subplots(
3     1, 2, figsize=(8, 6), sharex=True, sharey=True
4 )
5 fig.suptitle(f"Clustering Algorithm Comparison: Crescents", fontsize=16)
6 fte_colors = {
7     0: "#008fd5",
8     1: "#fc4f30",
9 }
10 # The k-means plot
11 kd_colors = [fte_colors[label] for label in kmeans1.labels_]
12 ax1.scatter(scaled_features1[:, 0], scaled_features1[:, 1], c=kd_colors)
13 ax1.set_title(
14     f"K-Means\nSilhouette: {silhouette_avg_mean[0]}", fontdict={"fontsize": 12}
15 )
16
17 # The ahc plot
18 ahc_colors = [fte_colors[label] for label in kmedoids1.labels_]
19 ax2.scatter(scaled_features1[:, 0], scaled_features1[:, 1], c=ahc_colors)
20 ax2.set_title(
21     f"K-Medoids\nSilhouette: {silhouette_avg_med[0]}", fontdict={"fontsize": 12}
22 )
23 plt.show()
```



5) Feature Selection Engineering

Feature Selection 1 - Filter Method (Removing Higher Correlated features)

```
In [34]: 1 def featSelectFilter(dfFeat1):
2         # Create correlation matrix
3         corr_matrix = corr.abs()
4
5         # Select upper triangle of correlation matrix
6         upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape), k=1).astype('bool'))
7
8         # Get features with correlation greater than 0.75
9         to_drop = [column for column in upper.columns if any(upper[column] > 0.75)]
10
11         print(color.BLUE + color.BOLD + 'Features removed since correlation is higher:' + color.END)
12         print(to_drop)
13
14         # Drop features
15         dfFeat1.drop(list(to_drop), axis=1, inplace=True)
16
17         #Size of the Dataset
18         print(color.BLUE + color.BOLD + "\nSize of Dataset:" + color.END)
19         print(dfFeat1.shape)
20
21         return dfFeat1
```

```
In [35]: 1 #Transform the non numerical data into numerical
2 def transformToNumerical(dfFeat1):
3     for column in dfFeat1.columns:
4         #If Column data type is Int i.e, numerical continue
5         if dfFeat1[column].dtype == 'int64' or dfFeat1[column].dtype == 'float64':
6             continue
7         #If Column data type is not Int, encode and transform to Numerical
8         dfFeat1[column] = LabelEncoder().fit_transform(dfFeat1[column].astype(str))
9
10    return dfFeat1
```

```
In [36]: 1 def minMaxScaler(dfFeat1):
2         # Using Min Max Scaler
3         min_max_scaler = MinMaxScaler()
4         min_max_scaled = min_max_scaler.fit_transform(dfFeat1)
5
6         # Creating new Data frame with the scaled value
7         FE1_Norm = pd.DataFrame(min_max_scaled, columns = dfFeat1.columns)
8
9         return FE1_Norm
```

```
In [37]: 1 #Split the entire dataset to Train and Test
2 def splitTrainTest(dfFeat1, dfFeat2):
3     #Splitting the dataset
4     X = dfFeat1.iloc[:, 0:dfFeat2.shape[1]]
5     X = X.drop(['RatingLevel', 'Rating'], axis=1)
6     Y = dfFeat1.iloc[:, -1]
7
8     #Split the data into 75% training and 25% testing
9     return train_test_split(X, Y, test_size = 0.25, random_state = 0)
```

```
In [38]: 1 #Dimensionality Reduction using PCA
2 def dimReductionPCA(X_train, X_test):
3     # Make an instance of the Model
4     pca = PCA(.95)
5
6     pca.fit(X_train)
7
8     X_train = pca.transform(X_train) #PCA transformation on Train Set
9     X_test = pca.transform(X_test) #PCA transformation on Test Set
10
11     #How much information (variance) attributed to each of the principal components
12     explained_variance = pca.explained_variance_ratio_
13     print(color.BLUE + color.BOLD + 'Variance attributed to each of the principal components:' + color.END)
14     print(explained_variance)
15
16     return X_train, X_test
```

```
In [39]: 1 #Handling the Dataset Imbalance Using Hybridization: SMOTE + Tomek Links
2 def handleClassImbalance(X_train, Y_train):
3
4     counter = Counter(Y_train) #Before Sampling, count of Y_train
5     print(color.BLUE + color.BOLD + 'Before Sampling:' + color.END)
6     print(counter)
7
8     #Oversampling the train dataset using SMOTE + Tomek
9     #To get better class clusters, Tomek Links are applied to oversampled minority class samples done by SMOTE
10    smtom = SMOTETomek(random_state=0)
11    X_train_smtom, y_train_smtom = smtom.fit_resample(X_train, Y_train) #Fit the resampled model
12
13    counter = Counter(y_train_smtom) #After Sampling, Count of y_train_smtom
14    print(color.BLUE + color.BOLD + 'After Sampling:' + color.END)
15    print(counter)
16
17    return X_train_smtom, y_train_smtom
```

```
In [40]: 1 # Train the model using Random Forest classifier
2 # This meta estimator fits a number of decision tree classifiers on sub-samples of the
3 # dataset and uses averaging to improve the predictive accuracy and control over-fitting
4
5 def randomForestModel(X_train_smtom, y_train_smtom):
6     #Number of trees given as '10' with criterion 'entropy' and seed for random generator is set as '0'
7     randomforest = RandomForestClassifier(n_estimators = 10, criterion = 'entropy', random_state = 0)
8     randomforest.fit(X_train_smtom, y_train_smtom)
9
10    return randomforest
```

```
In [41]: 1 #Plotting top features that help in predicting using the Random Forest Built-in Feature Importance
2
3 def importantFeatures(randomforest, dfFeat1):
4     #Determine the feature importance values
5     importances = randomforest.feature_importances_
6
7     #Create a dictionary with the importances values
8     important_features_dict = {}
9     for idx, val in enumerate(importances):
10         important_features_dict[idx] = val
11
12     #Sort the feature importances in descending order
13     important_features_list = sorted(important_features_dict,
14                                     key=important_features_dict.get,
15                                     reverse=True)[1:]
16
17     important_features = dfFeat1.columns[important_features_list]
18
19     #Visualize the top 6 feature importance using bar chart
20     feat_importances = pd.Series(importances[important_features_list], index=important_features)
21     feat_importances.nlargest(6).plot(kind='barh')
22     plt.xlabel('Feature Importance')
23     plt.title('Top Features With Higher Random Forest Feature Importance')
24     plt.show()
```


In [42]:

```

1  # Create Machine Learning models - Logistic regression and Decision tree to predict
2
3  def mlPredict(X_train_smtom, y_train_smtom, Y_test):
4      cv_dataFrames = []
5
6      # Prepare Machine Learning models - Logistic regression and Decision tree
7      models = []
8
9      #Parametric Supervised learning model based on probability
10     models.append(('Logistic Regression(LR)', LogisticRegression()))
11     #Non-Parametric Supervised learning model by learning simple decision rules inferred from the data features
12     models.append(('Decision Tree(CART)', DecisionTreeClassifier()))
13
14     results = []
15     mName = [] #List for collecting model names
16
17     #List of scoring metrics for comparison of models
18     scoring = ['accuracy', 'precision_weighted', 'recall_weighted', 'f1_weighted', 'roc_auc']
19
20     targ_names = ['Low', 'High'] #List of target values
21
22     for mName, model in models: #Looping through each of the models
23
24         #Split the dataset into '5' folds and Each fold is used once as a validation while the '5 - 1'
25         #remaining folds form the training set
26         #Shuffle is set to 'True' to shuffle the data before splitting into batches
27         #Random_state affects the ordering of the indices, which controls the randomness of each fold
28         kfold = model_selection.KFold(n_splits=5, shuffle=True, random_state=90210)
29
30         #Evaluate metrics by cross-validation
31         cv_res = model_selection.cross_validate(model, X_train_smtom, y_train_smtom, cv=kfold, scoring=scoring)
32
33         #Fit the model and predict the label of test set
34         ml = model.fit(X_train_smtom, y_train_smtom)
35         y_pred = ml.predict(X_test)
36
37         print(color.BLUE + color.BOLD + mName + color.END)
38
39         #Number of correct and incorrect predictions compared wih Actual class and Predicted class
40         cm = confusion_matrix(Y_test, y_pred)
41

```

```

42 TN = cm[0][0] #True Negative(Predicted No, Actual No, classifier is getting things right)
43 TP = cm[1][1] #True Positive(Predicted Yes, Actual Yes, classifier is getting things right)
44 FN = cm[1][0] #False Negative(Predicted No, Actual Yes, classifier is getting things wrong i.e, mislabel
45 FP = cm[0][1] #False Positive(Predicted Yes, Actual No, classifier is getting things wrong i.e, mislabel
46
47 print(color.BOLD + "\nConfusion Matrix:" + color.END)
48
49 column_names = ['Predicted Low', 'High']
50 row_names     = ['Actual Low', 'High']
51
52 cm_df = pd.DataFrame(cm, columns=column_names, index=row_names)
53
54 print(cm_df)
55
56 #Accuracy determines how often is classifier correct, (TP+TN)/Total
57 print(color.BOLD + "\nAccuracy:" + color.END)
58 print(round(accuracy_score(Y_test, y_pred) * 100, 2), "%")
59
60 #Return the list of scores calculated for each cv='10' folds, estimator object
61 #implementing 'fit' and n_jobs='-1' means using all processors
62 cross_val_lr = cross_val_score(estimator = model, X = X_train_smtom, y = y_train_smtom, cv = 10, n_jobs
63 print(color.BOLD + "\nCross Validation Accuracy:" + color.END)
64 print(round(cross_val_lr.mean() * 100 , 2), "%")
65
66 #Report showing the main classification metrics with the target names 'Yes' and 'No'
67 print(color.BOLD + "\nClassification Report:" + color.END)
68 print(classification_report(Y_test, y_pred, target_names=targ_names))
69
70 #Get False Positive Rates and True Postive rates for the Classifiers
71 #By roc_curve module by passing the test dataset and the predicted data through it
72 print(color.BOLD + "\nReceiver Operating Characteristic(ROC):" + color.END)
73 false_positive_rate1, true_positive_rate1, threshold1 = roc_curve(Y_test, y_pred)
74
75 #Ploting ROC Curves with False Positive Rate on X-axis and True Positive Rate on Y-axis
76 title = 'Receiver Operating Characteristic(ROC) - ' + mName
77 plt.subplots(1, figsize=(7,5))
78 plt.title(title)
79 plt.plot(false_positive_rate1, true_positive_rate1)
80 plt.plot([0, 1], ls="--")
81 plt.plot([0, 0], [1, 0] , c=".7"), plt.plot([1, 1] , c=".7")
82 plt.ylabel('True Positive Rate')
83 plt.xlabel('False Positive Rate')

```

```
84
85     results.append(cv_res) #Appending the cross validation metrics
86     mNameNames.append(mName) #Appending each of the model names
87
88     dataFrame = pd.DataFrame(cv_res) #Create data frame of cross validation results
89     dataFrame['model'] = mName #Add the model name to dataframe
90     cv_dataFrames.append(dataFrame) #Append each of the data frames
91
92     result = pd.concat(cv_dataFrames, ignore_index=True) #Concatenate the the dataframes object ingnoring in
93     return result
```

In [43]:

```

1  #Comparison of Performance of Logistic Regression and Decision tree models
2
3  def ml_ModelsComparison(result):
4      mlValues = []
5
6      #Iterating through result values and append the values of each models to mlValues[]
7      for model in list(set(result.model.values)):
8          m_dataFrame = result.loc[result.model == model]
9          mlValue = m_dataFrame.sample(n=30, replace=True)
10         mlValues.append(mlValue)
11
12         m_dataFrame = pd.concat(mlValues, ignore_index=True) #Concatenate the the dataframes object ingnoring index
13
14         #Massage a DataFrame into a format where identifier variable is 'model', variable column 'metrics'
15         #and value column 'values'
16         perf_results = pd.melt(m_dataFrame, id_vars=['model'], var_name='metrics', value_name='values')
17
18         tym_metrics = ['fit_time', 'score_time'] # Fit time Metrics
19
20         #Performance Metrics
21         perf_results_nofit = perf_results.loc[~perf_results['metrics'].isin(tym_metrics)] # Get dataframe without fit da
22         perf_results_nofit = perf_results_nofit.sort_values(by='values') #Sort the performance result on its values
23
24         #Visualization of Comparison of LR and CART Model using BoxPlot
25         plt.figure(figsize=(10, 7))
26         sns.set(font_scale=1)
27         g = sns.boxplot(x="model", y="values", hue="metrics", data=perf_results_nofit, palette="Set3")
28         plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
29         plt.title('Comparison of LR and CART Model by Classification Metric')
30
31         return perf_results_nofit, m_dataFrame

```

Feature Selection 2 - Chi-Square and Mutual Info

```
In [44]: 1 def select_featureschi(X_train, y_train):
2         fs = SelectKBest(score_func=chi2, k=5)
3         fs.fit_transform(X_train, y_train)
4         return fs
5
6 def selected_features_chi2(dfFeat2):
7     #Splitting the dataset
8     X = dfFeat2.iloc[:, 0:dfFeat2.shape[1]]
9     X = X.drop(['RatingLevel', 'Rating'], axis=1)
10    y = dfFeat2.iloc[:, -1]
11
12    fs = select_featureschi(X, y)
13
14    selected_features_chi2 = list(X.columns[fs.get_support(indices=True)])
15    return selected_features_chi2
```

```
In [45]: 1 def select_featuresinfo(X_train, y_train):
2         fs = SelectKBest(score_func=mutual_info_classif, k=5)
3         fs.fit_transform(X_train, y_train)
4         return fs
5
6 def selected_features_mutual_info(dfFeat2):
7     #Splitting the dataset
8     X = dfFeat2.iloc[:, 0:dfFeat2.shape[1]]
9     X = X.drop(['RatingLevel', 'Rating'], axis=1)
10    y = dfFeat2.iloc[:, -1]
11
12    fs = select_featuresinfo(X, y)
13
14    selected_features_mutual_info = list(X.columns[fs.get_support(indices=True)])
15    return selected_features_mutual_info
```

Calling the functions

```
In [65]: 1 #Feature Selection 1 - Filter Method(Removing Higher Correlated features)
2 dfFeature1 = df.copy() #Create a copy of dataset
3 dfFeature1 = featSelectFilter(dfFeature1)
```

Features removed since correlation is higher:

['Selling Price', 'Cost Price', 'Gross Income']

Size of Dataset:

(1000, 10)

```
In [66]: 1 #Feature Encoding (To Numerical)
2 dfFeature1 = transformToNumerical(dfFeature1)
3 dfFeature1.head(2)
```

Out[66]:

	City	Customer Type	Gender	Product Type	Unit Price	Quantity	Tax	Payment Type	Rating	RatingLevel
0	0	0	0	3	74.69	7	26.1415	2	9.1	0
1	2	1	0	0	15.28	5	3.8200	0	9.6	0

```
In [67]: 1 dfFeature1 = minMaxScaler(dfFeature1)
2 dfFeature1.head(2)
```

Out[67]:

	City	Customer Type	Gender	Product Type	Unit Price	Quantity	Tax	Payment Type	Rating	RatingLevel
0	0.0	0.0	0.0	0.6	0.718847	0.666667	0.521616	1.0	0.850000	0.0
1	1.0	1.0	0.0	0.0	0.057855	0.444444	0.067387	0.0	0.933333	0.0

```
In [68]: 1 #Splitting the Dataset to Train and Test
2 X_train, X_test, Y_train, Y_test = splitTrainTest(dfFeature1, dfFeature1)
```

```
In [69]: 1 #PCA Dimensionality Reduction
        2 X_train, X_test = dimReductionPCA(X_train, X_test)
```

Variance attributed to each of the principal components:

```
[0.22059842 0.19658315 0.14265224 0.14076454 0.12170646 0.0977056
 0.07657132]
```

```
In [70]: 1 #Handling Imbalanced Dataset
        2 X_train_smtom, y_train_smtom = handleClassImbalance(X_train, Y_train)
```


Before Sampling:

```
Counter({0.0: 510, 1.0: 240})
```

After Sampling:

```
Counter({1.0: 465, 0.0: 465})
```

```
In [71]: 1 #RandomForest Model for important features
        2 randomforest = randomForestModel(X_train_smtom, y_train_smtom)
        3 importantFeatures(randomforest, dfFeature1)
```



```
In [72]: 1 result = mlPredict(X_train_smtom, y_train_smtom, Y_test)
```

Logistic Regression(LR)

Confusion Matrix:

	Predicted Low	High
Actual Low	88	61
High	50	51

Accuracy:

55.6 %

Cross Validation Accuracy:

52.26 %

Classification Report:

	precision	recall	f1-score	support
Low	0.64	0.59	0.61	149
High	0.46	0.50	0.48	101
accuracy			0.56	250
macro avg	0.55	0.55	0.55	250
weighted avg	0.56	0.56	0.56	250

Receiver Operating Characteristic(ROC):

Decision Tree(CART)

Confusion Matrix:

	Predicted Low	High
Actual Low	102	47
High	63	38

Accuracy:

56.0 %


Cross Validation Accuracy:


66.45 %

Classification Report:

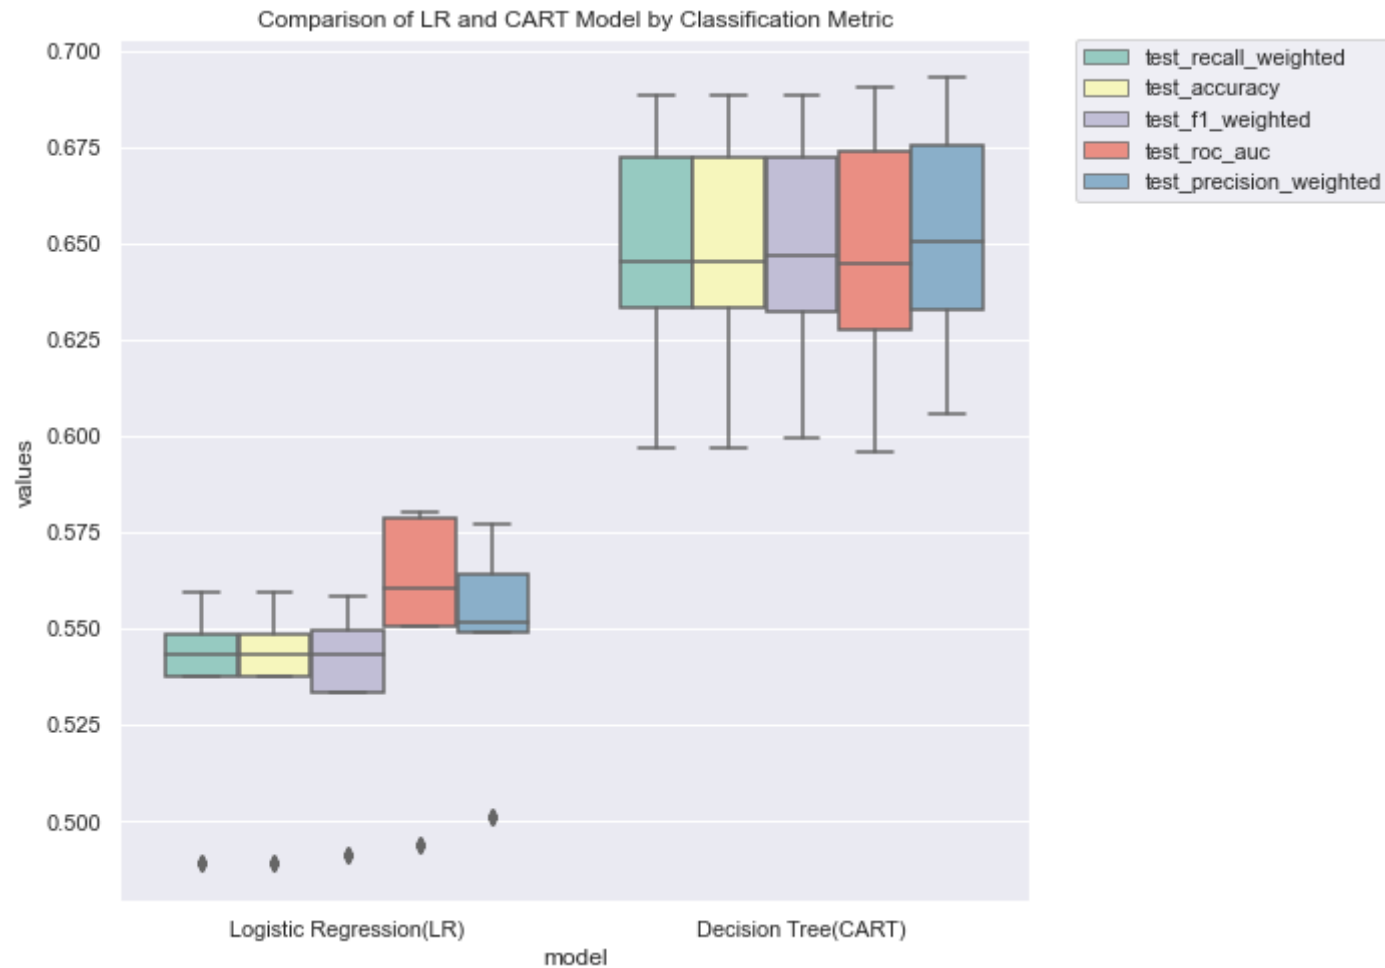
	precision	recall	f1-score	support
Low	0.62	0.68	0.65	149
High	0.45	0.38	0.41	101
accuracy			0.56	250
macro avg	0.53	0.53	0.53	250
weighted avg	0.55	0.56	0.55	250

Receiver Operating Characteristic(ROC):





```
In [73]: 1 #Machine Learning models - Logistic regression and Decision tree to predict attrition and Comparison of Performance  
2 perf_results_nofit, m_dataFrame = ml_ModelsComparison(result)
```



Comparison of Performance metrics

```
In [74]: 1 metricValues = list(set(perf_results_nofit.metrics.values))
2 #aggregate metric values with standard deviation and mean
3 m_dataFrame.groupby(['model'])[metricValues].agg([np.std, np.mean])
```

Out[74]:

	test_roc_auc		test_precision_weighted		test_recall_weighted		test_f1_weighted		test_accuracy	
	std	mean	std	mean	std	mean	std	mean	std	mean
model										
Decision Tree(CART)	0.030885	0.652509	0.028650	0.656883	0.028614	0.653047	0.028424	0.653472	0.028614	0.653047
Logistic Regression(LR)	0.028116	0.557448	0.022785	0.551660	0.021233	0.539247	0.020781	0.538785	0.021233	0.539247

```
In [95]: 1 dfFeature2 = df
2 dfFeature2 = transformToNumerical(dfFeature2)
3 featSelect2 = selected_features_chi2(dfFeature2) + selected_features_mutual_info(dfFeature2)
4 featSelect2 = list(set(featSelect2))
5 dfFeature2 = dfFeature2[featSelect2]
6 dfFeature2.head(2)
```

Out[95]:

	Tax	Selling Price	Cost Price	Gross Income	Unit Price	City
0	26.1415	548.9715	522.83	26.1415	74.69	0
1	3.8200	80.2200	76.40	3.8200	15.28	2

```
In [96]: 1 dfFeature2 = minMaxScaler(dfFeature2)
        2 dfFeature2.head(2)
```

Out[96]:

	Tax	Selling Price	Cost Price	Gross Income	Unit Price	City
0	0.521616	0.521616	0.521616	0.521616	0.718847	0.0
1	0.067387	0.067387	0.067387	0.067387	0.057855	1.0

```
In [97]: 1 #Splitting the Dataset to Train and Test
        2 X = dfFeature2.iloc[:, 0:dfFeature2.shape[1]]
        3 Y = df.iloc[:, -1]
        4
        5 #Split the data into 75% training and 25% testing
        6 X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.25, random_state = 0)
```

```
In [98]: 1 #PCA Dimensionality Reduction
        2 X_train, X_test = dimReductionPCA(X_train, X_test)
```

Variance attributed to each of the principal components:
[0.56511135 0.34336016 0.09152849]

```
In [99]: 1 #Handling Imbalanced Dataset
        2 X_train_smtom, y_train_smtom = handleClassImbalance(X_train, Y_train)
```

Before Sampling:


Counter({0: 510, 1: 240})

After Sampling:

Counter({1: 420, 0: 420})

In [100]:

```
1 #RandomForest Model for important features
2 randomforest = RandomForestModel(X_train_smtom, y_train_smtom)
3 importantFeatures(randomforest, dfFeature2)
```



```
In [101]: 1 #Implement machine Learning models to predict and gets the result of model and its performance metric values
          2 result = mlPredict(X_train_smtom, y_train_smtom, Y_test)
```

Logistic Regression(LR)

Confusion Matrix:

	Predicted Low	High
Actual Low	92	57
High	52	49

Accuracy:

56.4 %

Cross Validation Accuracy:

53.1 %

Classification Report:

	precision	recall	f1-score	support
Low	0.64	0.62	0.63	149
High	0.46	0.49	0.47	101
accuracy			0.56	250
macro avg	0.55	0.55	0.55	250
weighted avg	0.57	0.56	0.57	250

Receiver Operating Characteristic(ROC):

Decision Tree(CART)

Confusion Matrix:

	Predicted Low	High
Actual Low	95	54
High	63	38

Accuracy:


53.2 %


Cross Validation Accuracy:

71.43 %

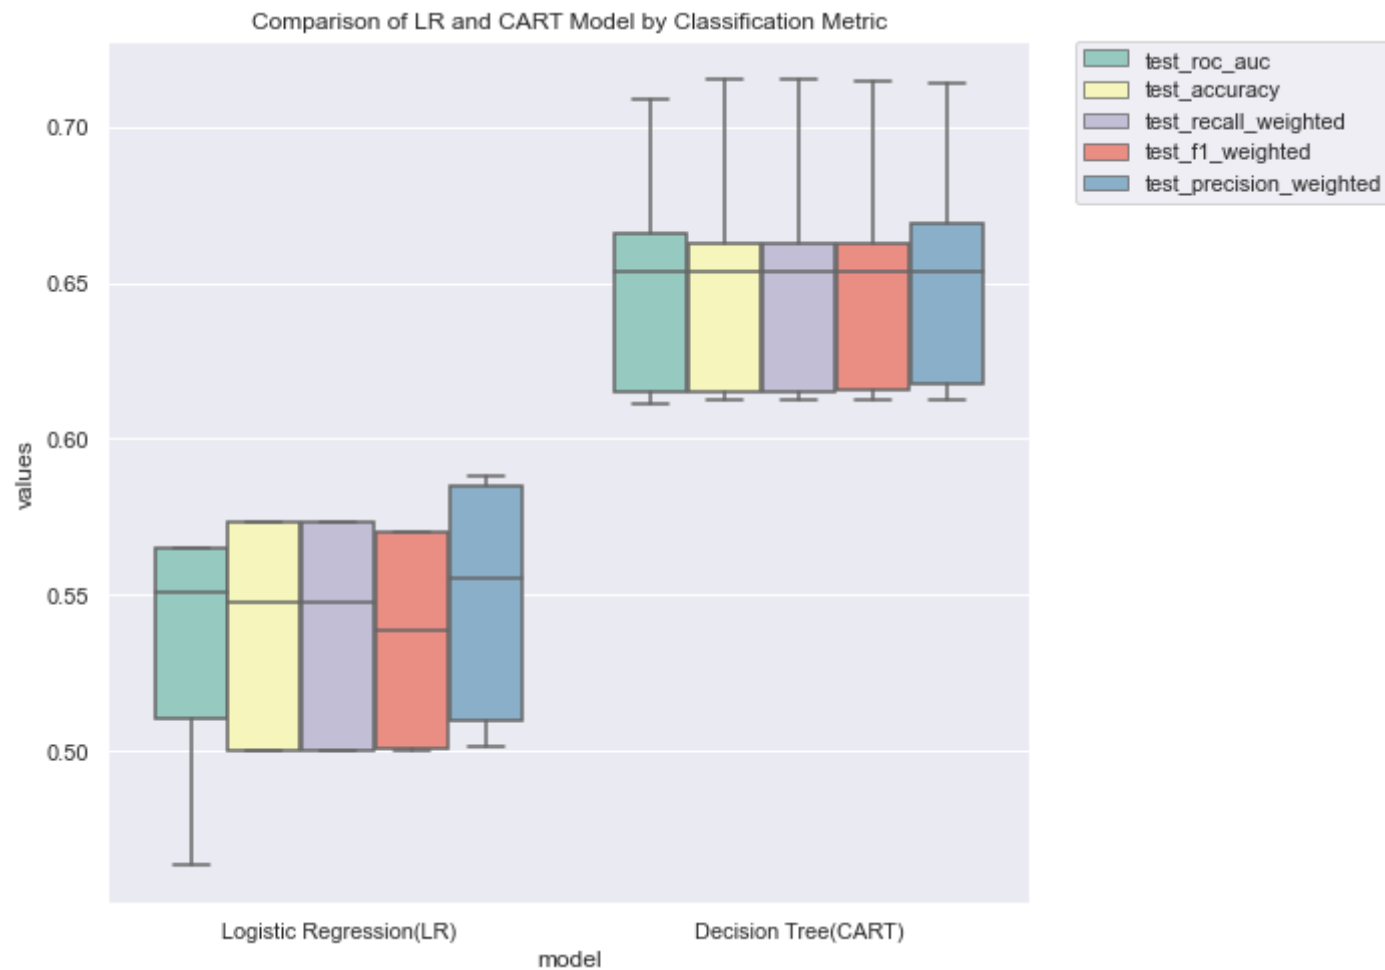
Classification Report:

	precision	recall	f1-score	support
Low	0.60	0.64	0.62	149
High	0.41	0.38	0.39	101
accuracy			0.53	250
macro avg	0.51	0.51	0.51	250
weighted avg	0.53	0.53	0.53	250

Receiver Operating Characteristic(ROC):



```
In [93]: 1 #Machine Learning models - Logistic regression and Decision tree to predict attrition and Comparison of Performance  
2 perf_results_nofit, m_dataFrame = ml_ModelsComparison(result)
```



In [102]:

```

1 #Comparison of Performance metrics
2 metricValues = list(set(perf_results_nofit.metrics.values))
3 #aggregate metric values with standard deviation and mean
4 m_dataframe.groupby(['model'])[metricValues].agg([np.std, np.mean])

```

Out[102]:

	test_roc_auc		test_precision_weighted		test_recall_weighted		test_f1_weighted		test_accuracy	
	std	mean	std	mean	std	mean	std	mean	std	mean
model										
Decision Tree(CART)	0.036726	0.653031	0.037728	0.655696	0.038920	0.653889	0.038575	0.653844	0.038920	0.653889
Logistic Regression(LR)	0.039884	0.532562	0.036826	0.550070	0.031268	0.538111	0.030115	0.534811	0.031268	0.538111